


Java Parallel Proccessing Framework



***An Open Source alternative to
grid computing***

by Laurent Cohen

Agenda

- Why JPPF
- Main features
- General architecture
- How to use it with my own Java program
- How to get rid of deployment hassles
- Optimized for performance
- Robustness and scalability
- A word about security
- Administration and monitoring
- And tomorrow...

The origins of the framework

- Tackling AI problems
 - Neural networks
 - Genetic algorithms
- Very resource hungry
- Why not better utilize existing hardware?
- Need for parallelization
- Issues with existing frameworks
 - Heavy
 - Complex to setup and use
 - Not platform independent
 - Not Open-Source

How it came to life

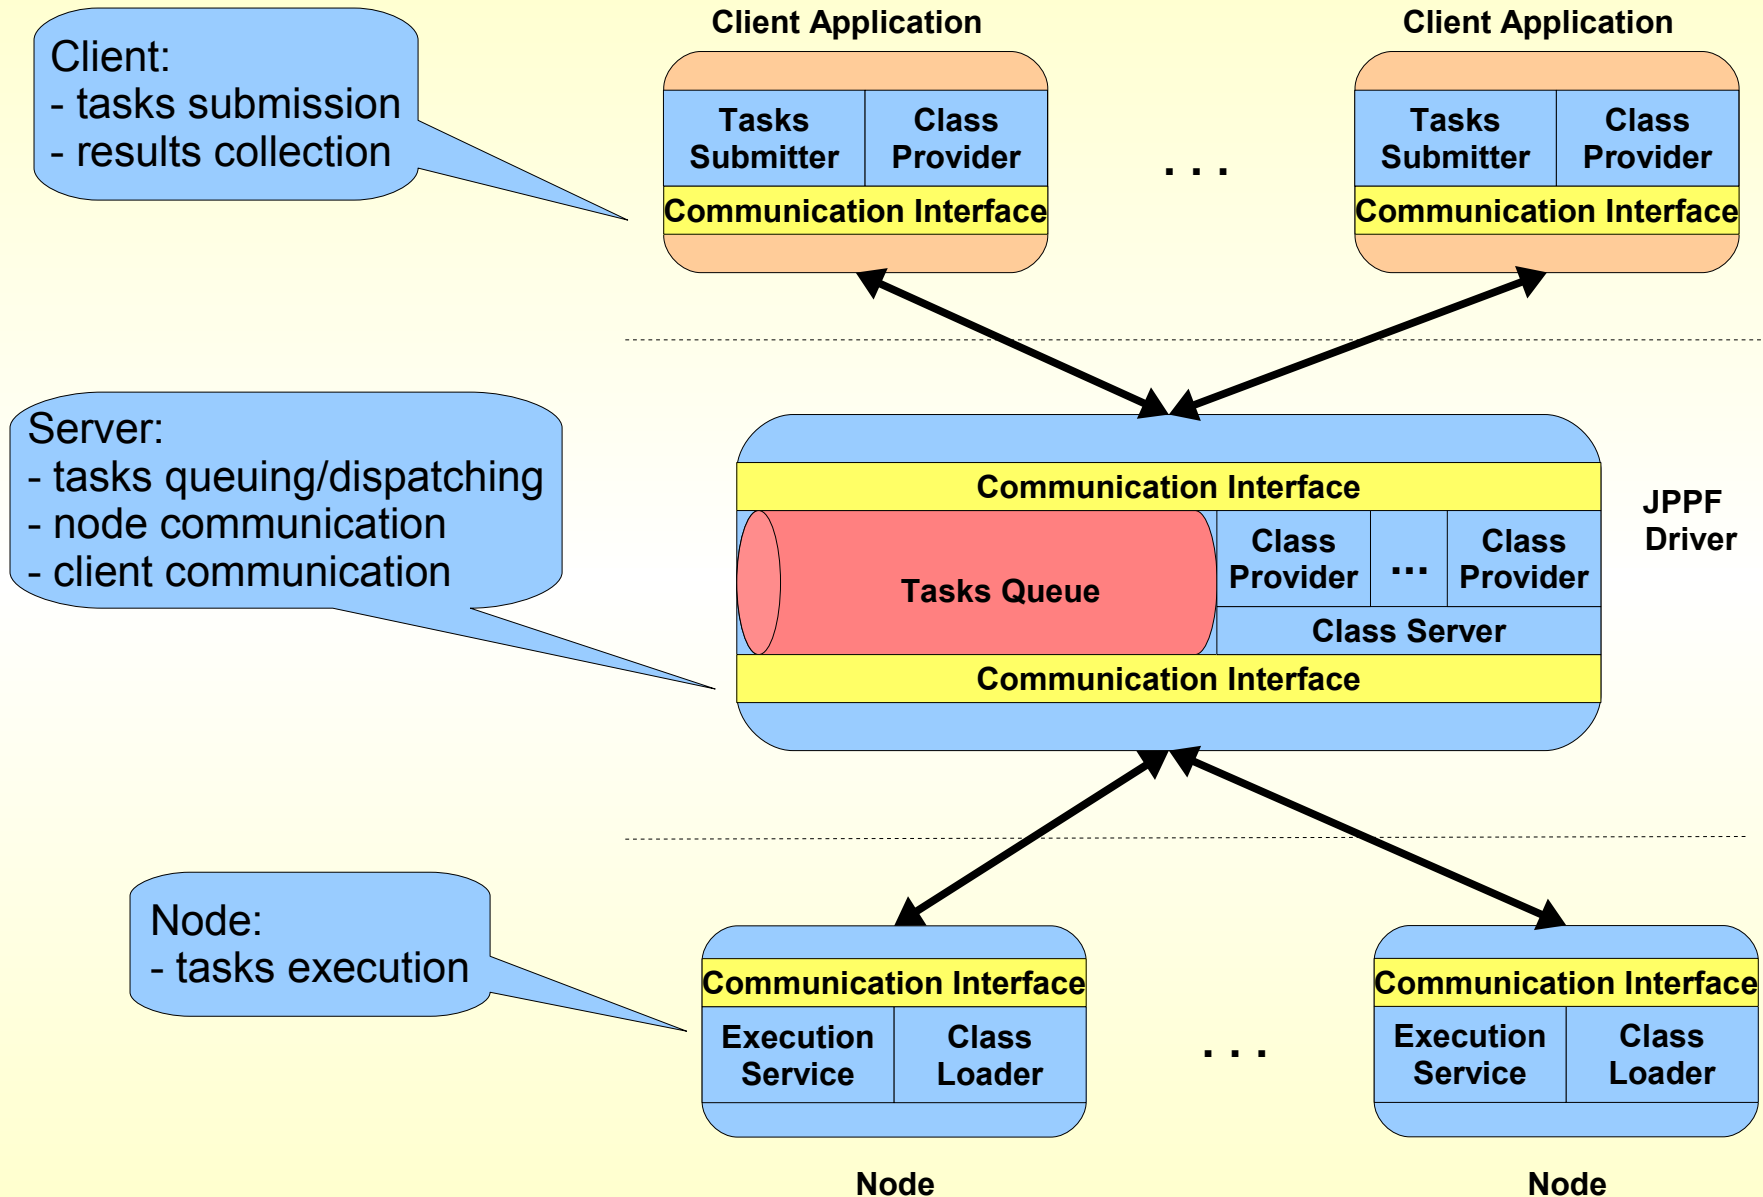
- Excitement for the project
 - parallel computing is a great challenge
 - Open-Source adventure (thanks SF.net)
- Need to parallelize a variety of applications
- Need to work on heterogenous environments
- Need for performance and scalability
- Confidence in software engineering skills
- Need for ubiquitous, platform-independant language
 - Java is a great choice
 - plenty of professional and personal experience

Main Features

What JPPF does

- API to execute independent tasks in parallel
- Scalability to many thousands of nodes
- Platform-independance for nodes and clients
- Execute your code without need for deployment
- Code executed in a secure environment
- Fail-over and recovery
- Graphical monitoring of the framework health
- Remote administration

General Architecture



A glimpse of the API: defining an atomic task

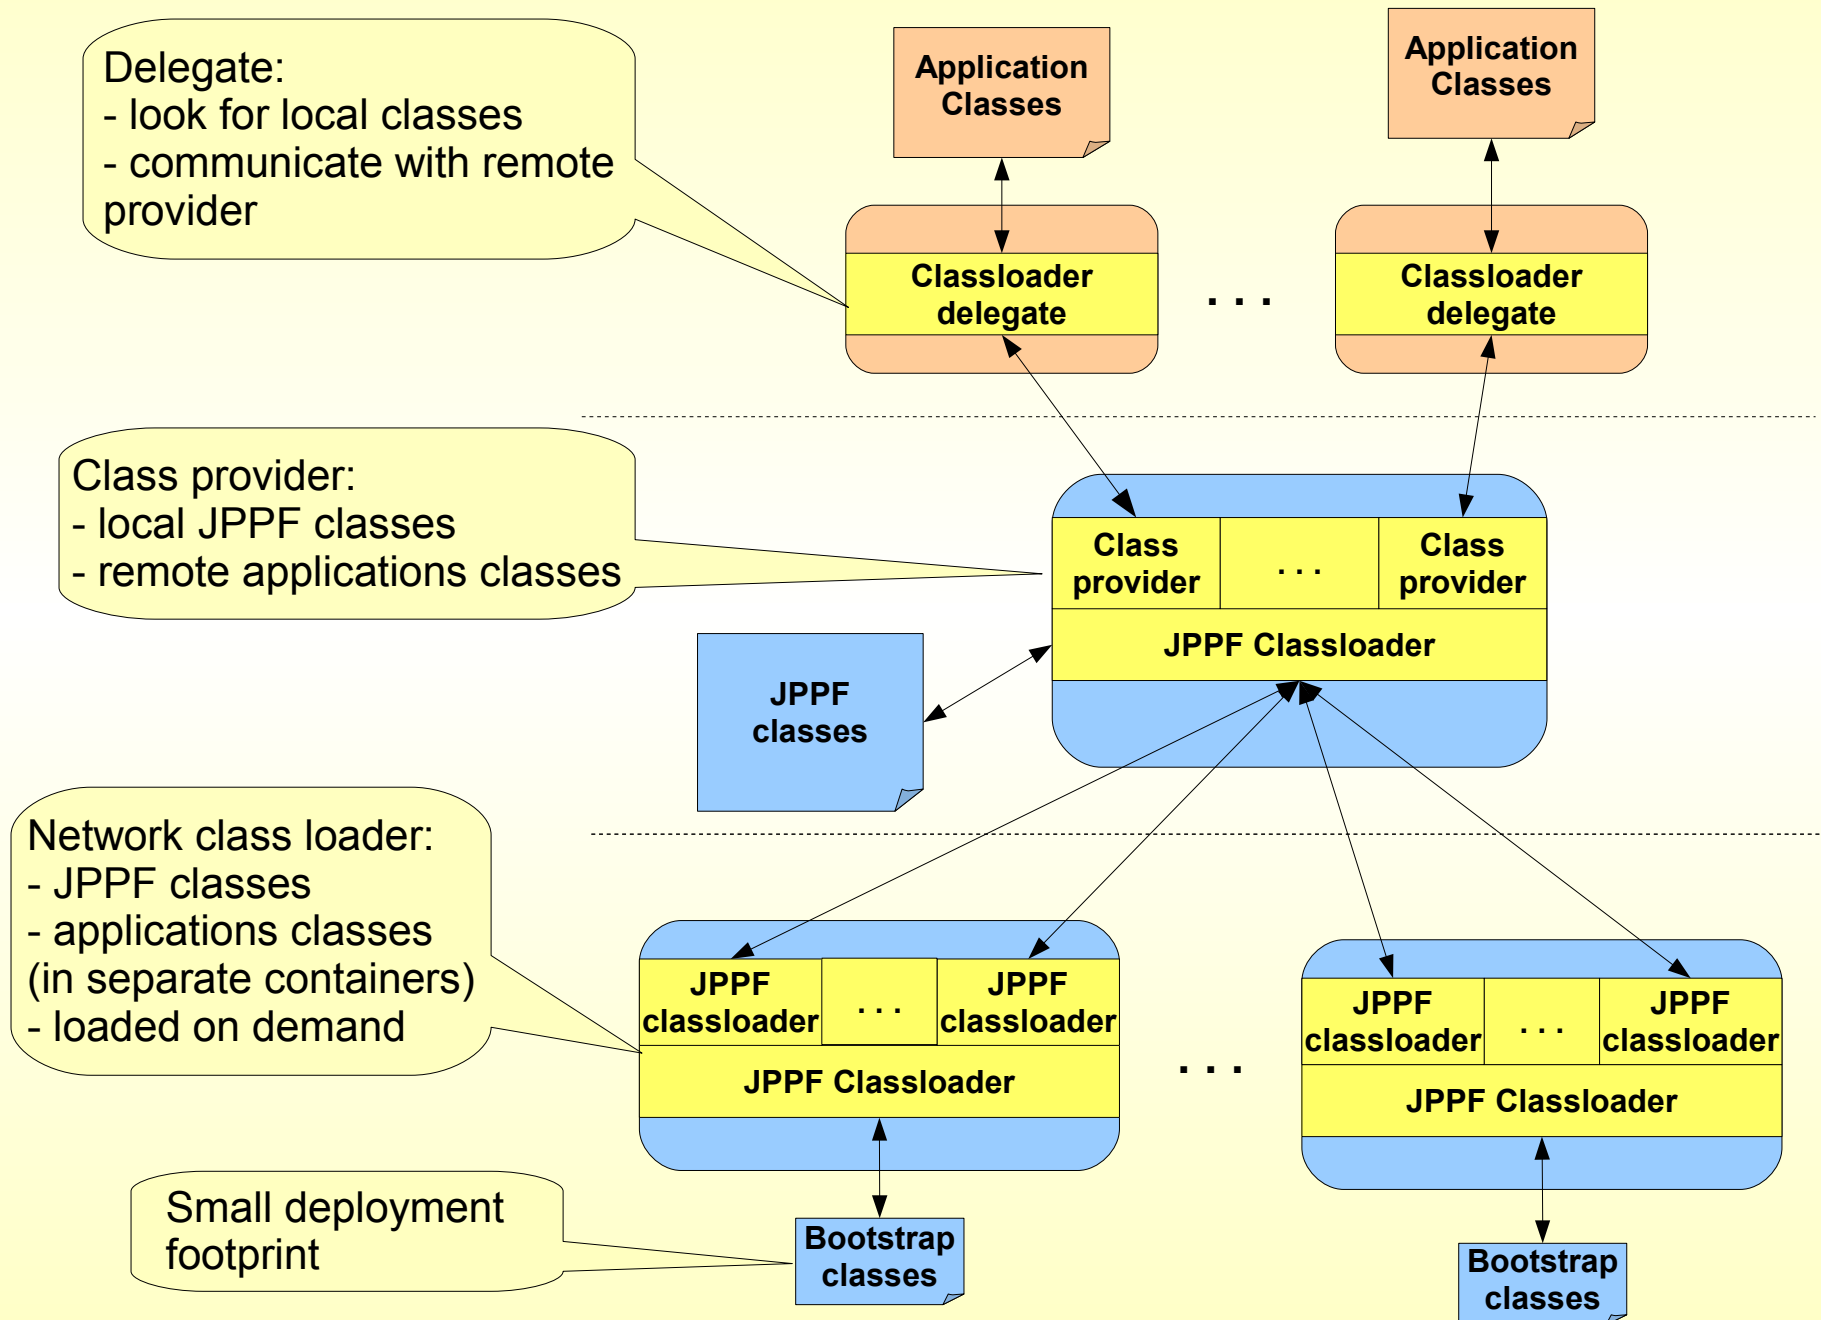
```
/**
 * This class performs my parallelized code.
 * @author Laurent Cohen
 */
public class MyTask extends JPPFTask
{
    /**
     * This is where it happens.
     * @see java.lang.Runnable#run()
     */
    public void run()
    {
        // here it begins
        ... my code ...
        // here it ends
        setResult(theComputationResult);
    }
}
```

A glimpse of the API: submitting tasks for execution

```
/**
 * This class submits atomic tasks for execution.
 */
public class TaskSubmitter
{
    /**
     * Submit the tasks and wait for their execution.
     * @throws Exception if any exception is raised.
     */
    public void submit() throws Exception
    {
        JPPFClient jppfClient = new JPPFClient();
        List<JPPFTask> taskList = new ArrayList<JPPFTask>();
        taskList.add(new MyTask());
        taskList.add(new MyTask());
        jppfClient.submit(taskList, null);
        System.out.println("The result of the first task is: "
            + taskList.get(0).getResult());
    }
}
```

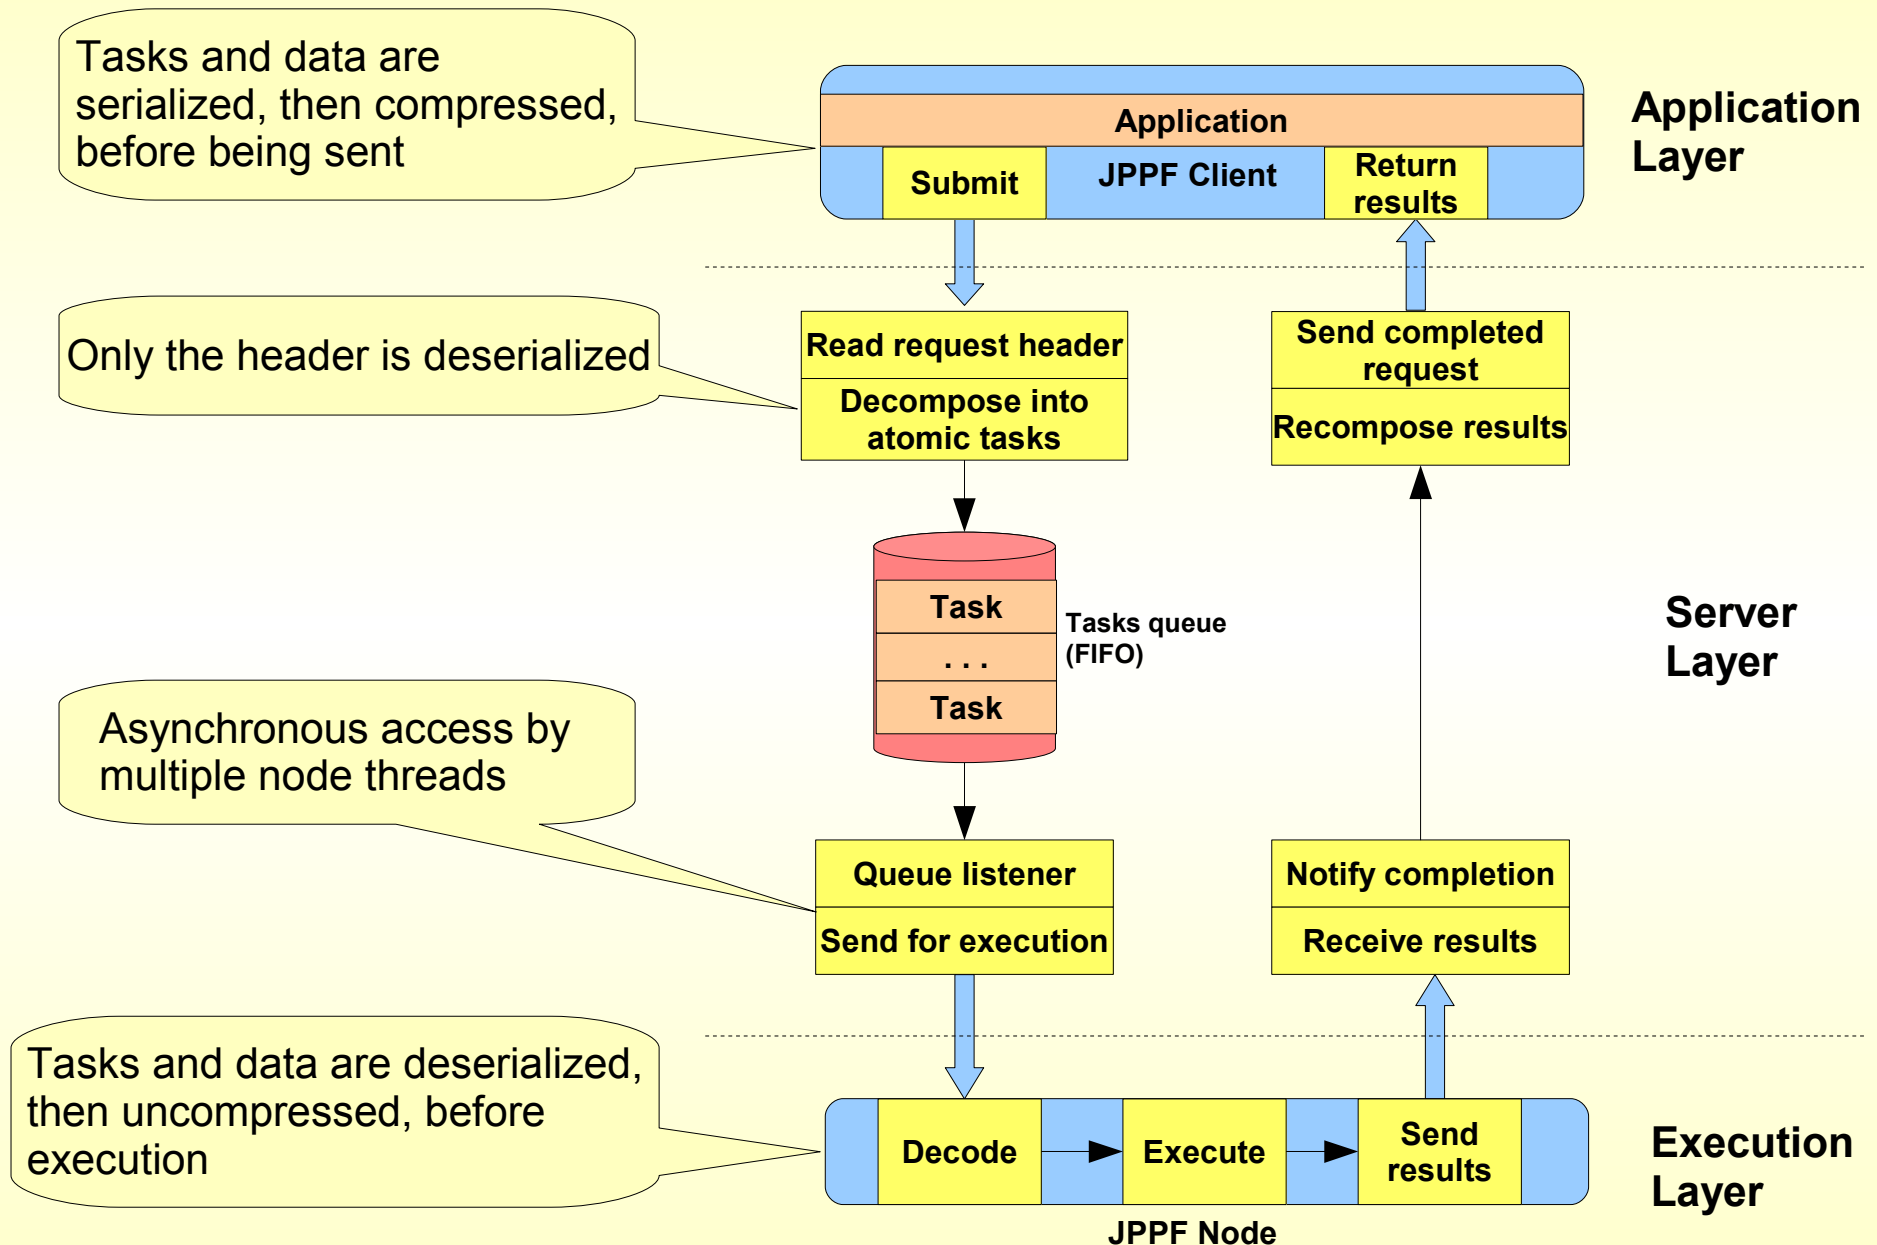

How to get rid of deployment hassles

1/2



How do I benefit from dynamic code loading?

- No need to deploy classes and libraries
 - The API takes care of it
- Code updates automatically accounted for
- Security of separate containers
 - Other applications don't see my code
 - I can't see other applications code
- Framework updates automatically propagated
 - The nodes download their own code



How JPPF deals with major performance sinks

- The server is an obvious bottleneck
 - Doesn't need to see the code as such
 - Asynchronous queue removes management overhead
 - Server code profiled to minimize the CPU utilization
- Network transport optimization
 - Compression/decompression of tasks code, data and results on the client/node side
 - Asynchronous communication limits network traffic
 - TCP sockets rather than high-level object broker

Robustness and scalability

What it takes to have a high quality of service

- Failover and recovery
 - Nodes and clients reconnect automatically
 - Clients and server automatically resubmit tasks
 - Configurable to adapt to environment
- Current architecture limits
 - Single server architecture
 - Number of nodes limited by available TCP ports
 - 2 connections per node, so only 30,000 nodes max
- Server hardware a major factor
 - Multi-cpu highly recommended

A word about security

- Nodes can't do everything
 - File system access restrictions
 - Network connection to the JPPF server only
 - Limited access to system environment
 - Prevented from halting the JVM
 - Rights customizable through a policy file
- My code confidentiality is guaranteed
 - Runs in an isolated container
 - The server doesn't see it
- Administration requires authentication
 - No password sent in clear over the network

Administration and monitoring

Keeping control over the framework

- Capability to stop / pause / restart the server
 - Allow maintenance with minimal interruption
 - Handle overload situations
- Execution and server health data collection
 - Latest statistics
 - Variety of real-time charts
 - Fully customizable charting capabilities
- Graphical tool built on top of the client API
 - Acts as a client with special capabilities
 - Easy to plugin different kinds of tools

And tomorrow...

Project vision

- Multi-server architecture
 - Extend the scalability to millions of nodes
- Automate the server behavior
 - Configurable alert conditions
 - Automatic responses to alerts
- Pay per use grids: accounting for the time spent
- JPPF@Home
 - Node as screen saver, applet, Java Web Start app.
 - Keep the Open-Source spirits up!